

Kunal Saini

📍 Sonipat, Haryana - 131001, India

✉ 2112sainikunal@gmail.com

📞 +91 7206140828

🌐 LinkedIn

🐙 GitHub

Summary

AI Engineer specializing in Agentic AI, Enterprise Knowledge Systems, Retrieval Infrastructure, AI Safety, and Production AI Platforms. Architected scalable multi-agent frameworks, hybrid retrieval pipelines, and human-in-the-loop observability operations for enterprise environments.

Skills

- **Agentic AI & Orchestration:** LangGraph, Multi-Agent Systems, MCP, Tool Calling
- **Evaluation & Observability:** RAGAS, DeepEval, Langfuse, LangSmith, Human-in-the-Loop
- **Infrastructure & Deployment:** Docker, AWS, CI/CD, GitHub Actions, Vercel, Render
- **Retrieval Systems:** GraphRAG, RAG, Knowledge Graphs, Vector Search, Hybrid Search
- **Backend & Data:** Python, FastAPI, Neo4j, Qdrant, PostgreSQL, REST APIs

Engineering Projects

ANDROMEDA — Enterprise Agent Platform

[🔗 Live](#) [🐙 GitHub](#)

- **Architected** a scalable LangGraph-based enterprise agent platform integrating MCP tool ecosystems and retrieval workflows for production AI operations.
- **Orchestrated** multi-agent systems and Tool Calling capabilities via FastAPI, automating complex enterprise task execution pipelines.
- **Established** a comprehensive observability suite using DeepEval, RAGAS, LangSmith, and Langfuse to monitor LLMOps, track latency, and improve system reliability.
- **Integrated** policy enforcement engines and human-in-the-loop escalation protocols to guarantee deterministic and safe agent behavior.

CENTAURUS — Enterprise Knowledge Worker Platform

[🔗 Live](#) [🐙 GitHub](#)

- **Engineered** an enterprise knowledge operating system combining GraphRAG and hybrid search to power intelligent publishing infrastructure.
- **Constructed** knowledge graphs and identity resolution pipelines using Neo4j and Qdrant, accelerating context-aware vector search across complex data.
- **Implemented** multi-agent retrieval workflows leveraging LangGraph and MCP, significantly increasing data extraction accuracy and synthesis speed.
- **Optimized** embedding models and retrieval accuracy by designing continuous Human-in-the-Loop governance and feedback loops.

OSKAR — AI Safety & Moderation Platform

[🔗 Live](#) [🐙 GitHub](#)

- **Developed** a production-ready AI safety platform combining graph reasoning and vector retrieval to enforce contextual moderation policies at scale.
- **Designed** a moderation engine utilizing Graph Neural Networks (GraphSAGE), Transformer NLP, and GraphRAG against Neo4j and Qdrant databases.
- **Automated** the routing of ambiguous edge cases to human escalation workflows by applying Shannon entropy-based uncertainty and Bayesian trust models.
- **Enabled** complete transparency in automated safety operations by mapping explainable moderation decisions directly back to source graph structures.

Education

Bachelor of Technology (Computer Science Engineering – AI & ML)

Panipat Institute of Engineering and Technology

Aug 2022 – May 2026

Panipat, Haryana - 132102